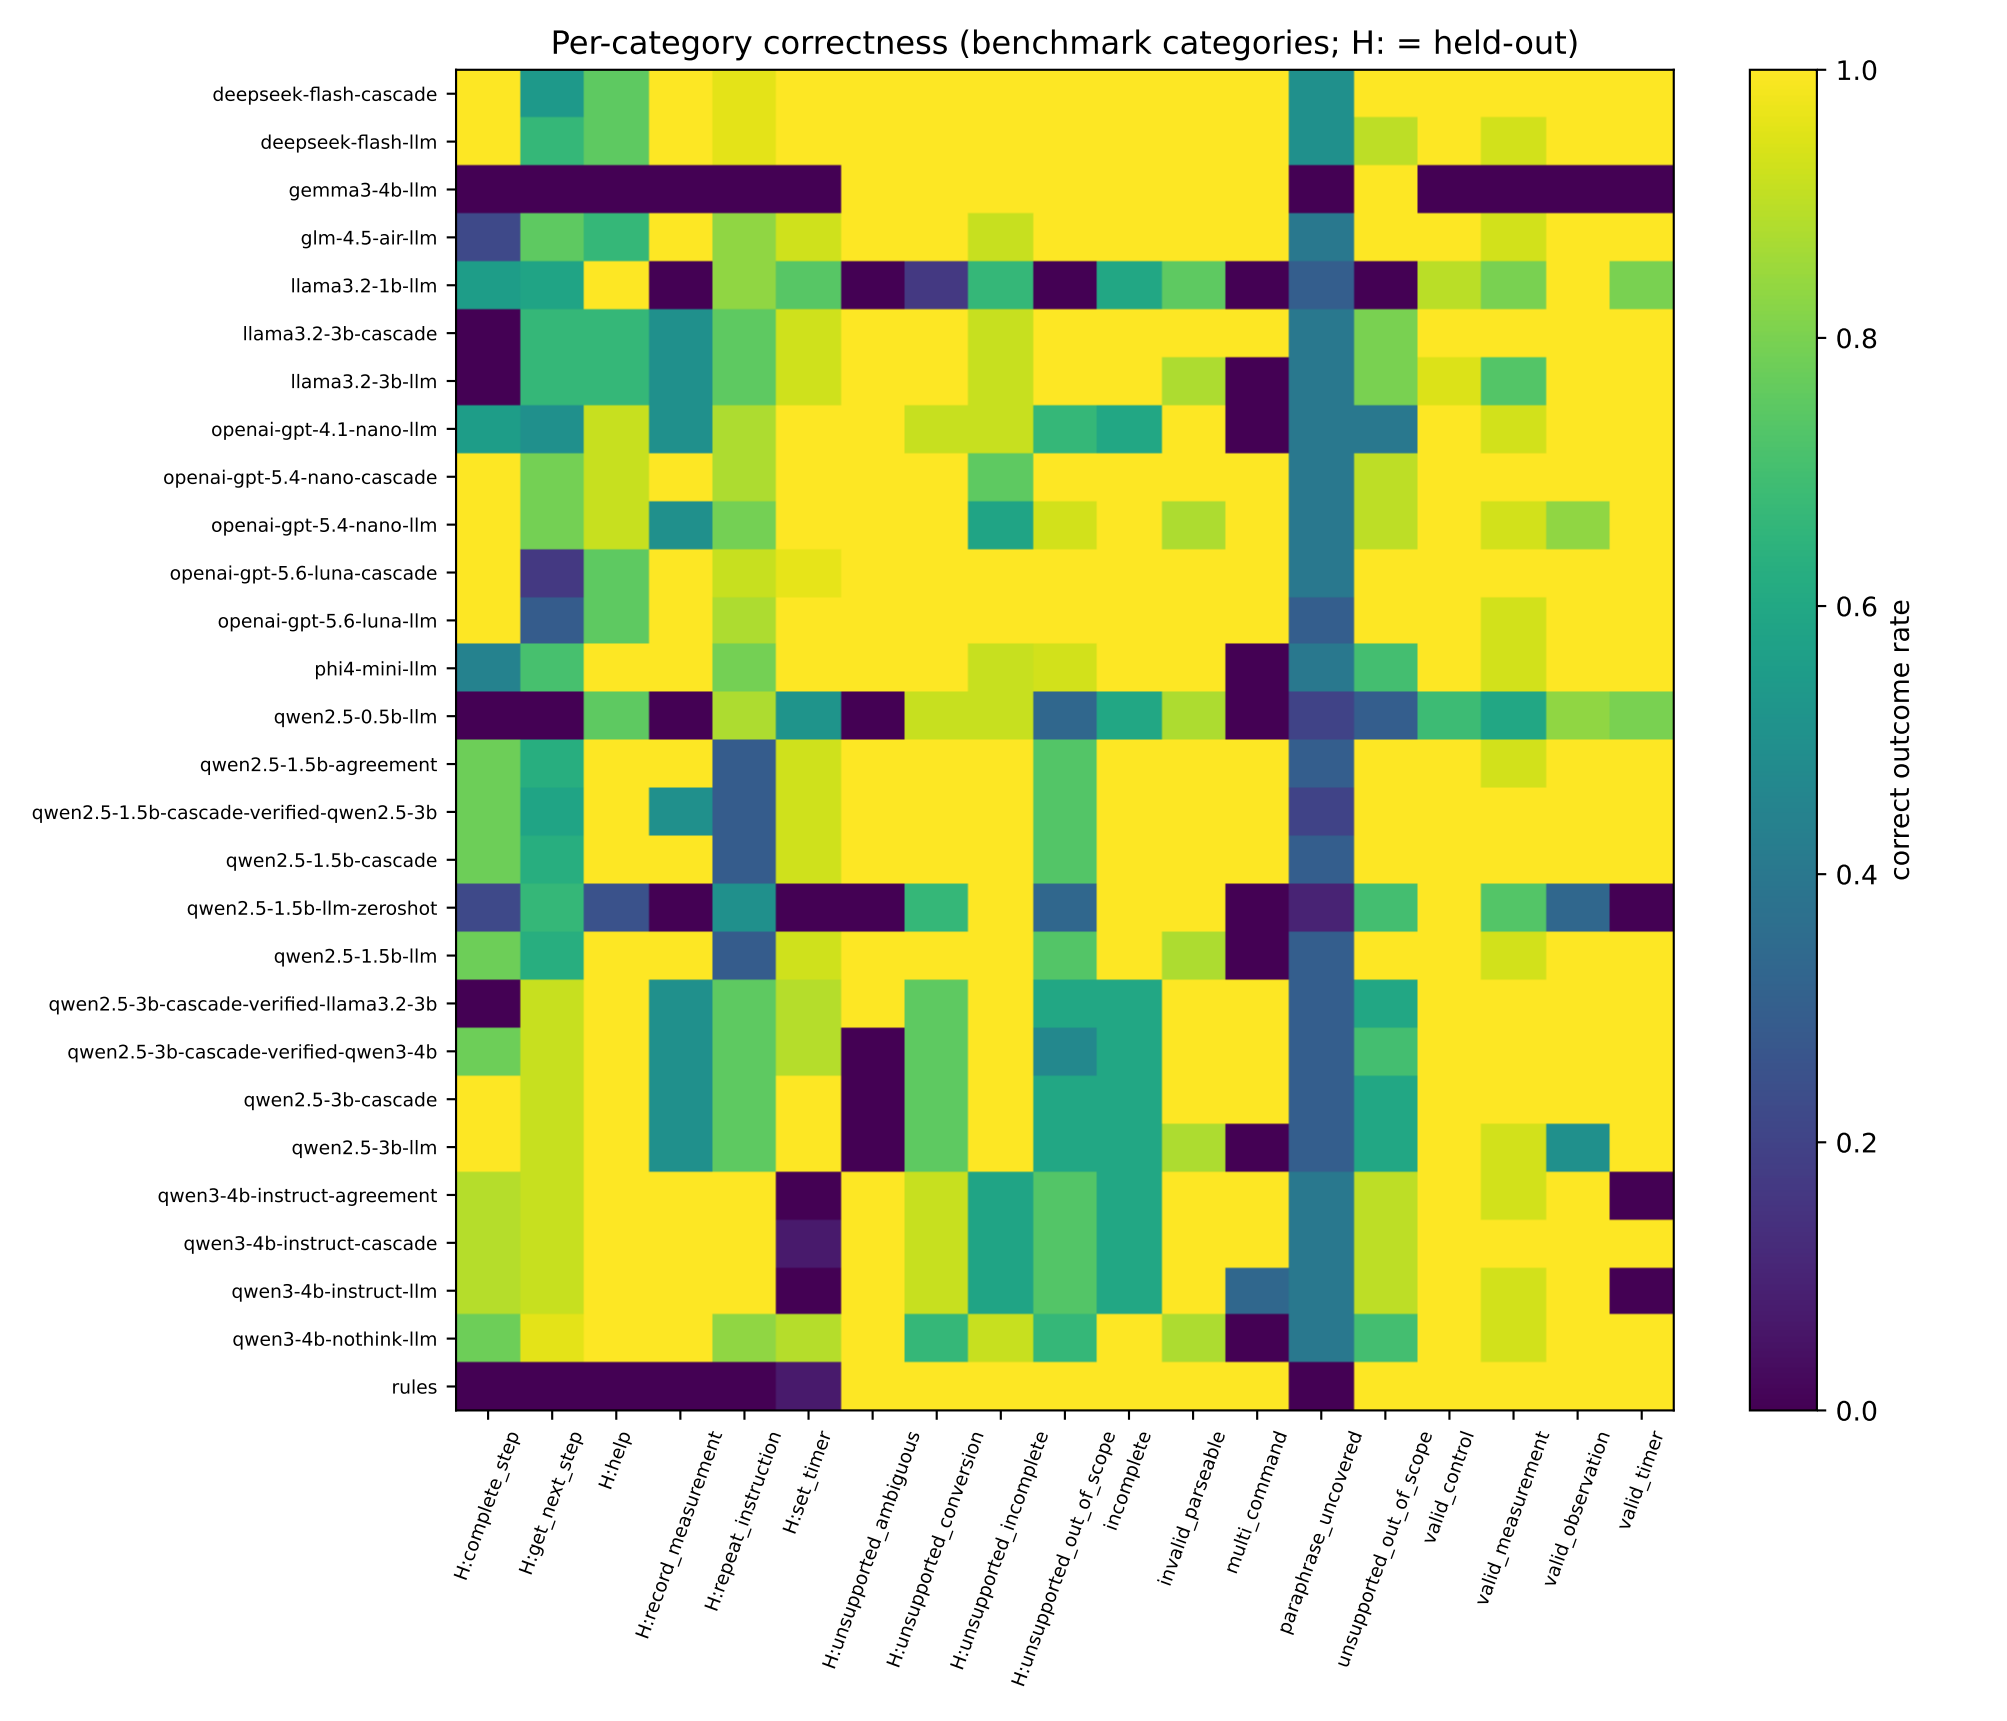
Per-category correctness (benchmark categories; H: = held-out)
deepseek-flash-cascade
deepseek-flash-llm
gemma3-4b-llm
glm-4.5-air-llm
llama3.2-1b-llm
llama3.2-3b-cascade
llama3.2-3b-llm
openai-gpt-4.1-nano-llm
openai-gpt-5.4-nano-cascade
openai-gpt-5.4-nano-llm
openai-gpt-5.6-luna-cascade
openai-gpt-5.6-luna-llm
phi4-mini-llm
qwen2.5-0.5b-llm
qwen2.5-1.5b-agreement
qwen2.5-1.5b-cascade-verified-qwen2.5-3b
qwen2.5-1.5b-cascade
qwen2.5-1.5b-llm-zeroshot
qwen2.5-1.5b-llm
qwen2.5-3b-cascade-verified-llama3.2-3b
qwen2.5-3b-cascade-verified-qwen3-4b
qwen2.5-3b-cascade
qwen2.5-3b-llm
qwen3-4b-instruct-agreement
qwen3-4b-instruct-cascade
qwen3-4b-instruct-llm
qwen3-4b-nothink-llm
rules
H:complete_step
H:get_next_step
H:help
H:record_measurement
H:repeat_instruction
H:set_timer
H:unsupported_ambiguous
H:unsupported_conversion
H:unsupported_incomplete
H:unsupported_out_of_scope
incomplete
invalid_parseable
multi_command
paraphrase_uncovered
unsupported_out_of_scope
valid_control
valid_measurement
valid_observation
valid_timer
correct outcome rate
1.0
0.8
0.6
0.4
0.2
0.0