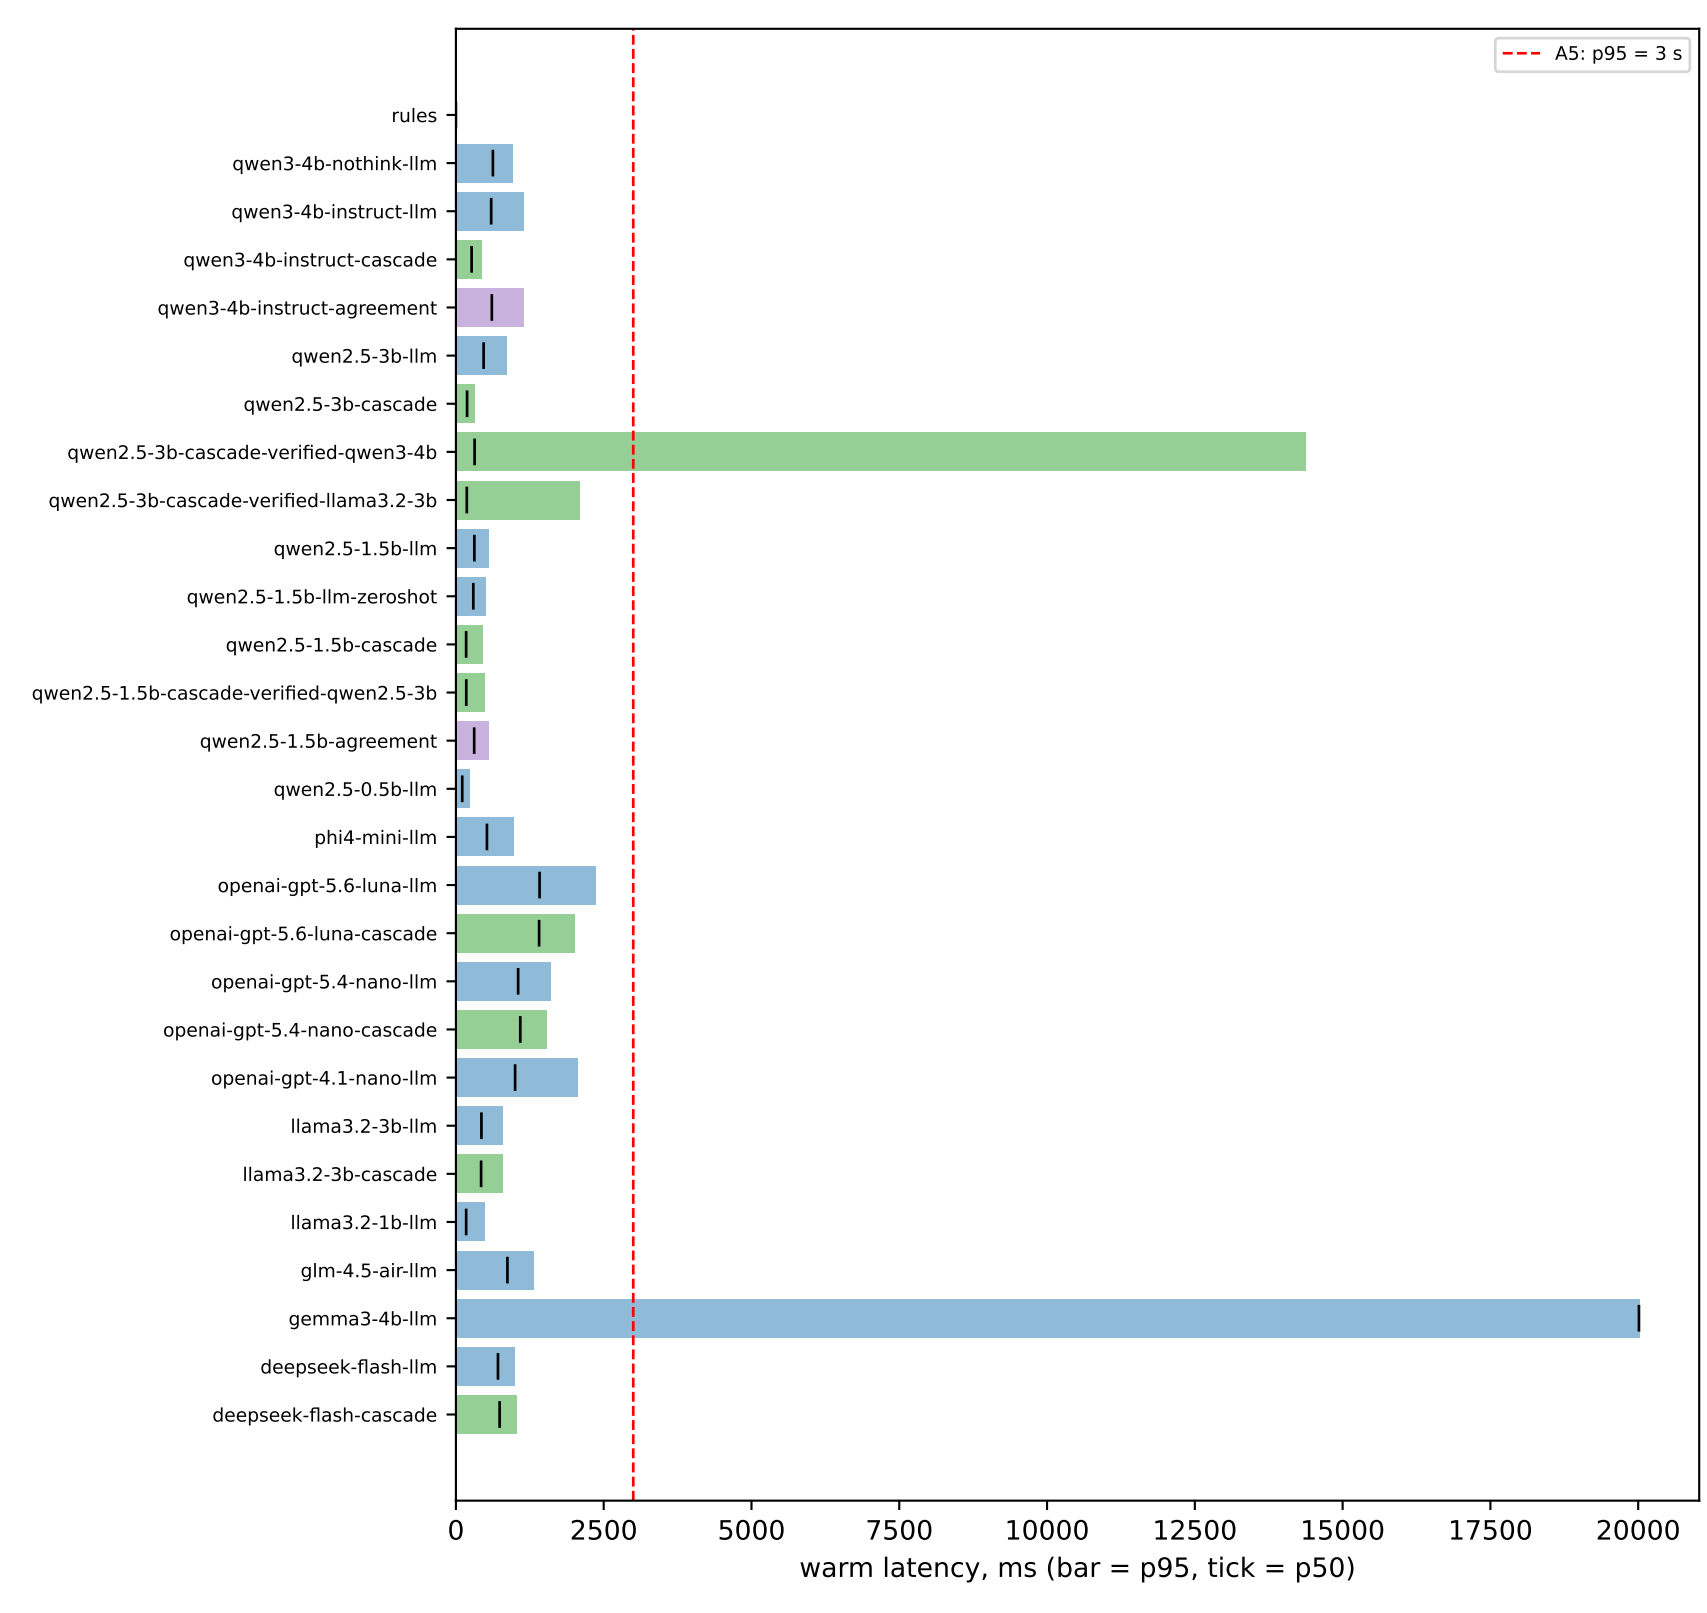
A5: p95 = 3 s
rules
qwen3-4b-nothink-llm
qwen3-4b-instruct-llm
qwen3-4b-instruct-cascade
qwen3-4b-instruct-agreement
qwen2.5-3b-llm
qwen2.5-3b-cascade
qwen2.5-3b-cascade-verified-qwen3-4b
qwen2.5-3b-cascade-verified-llama3.2-3b
qwen2.5-1.5b-llm
qwen2.5-1.5b-llm-zeroshot
qwen2.5-1.5b-cascade
qwen2.5-1.5b-cascade-verified-qwen2.5-3b
qwen2.5-1.5b-agreement
qwen2.5-0.5b-llm
phi4-mini-llm
openai-gpt-5.6-luna-llm
openai-gpt-5.6-luna-cascade
openai-gpt-5.4-nano-llm
openai-gpt-5.4-nano-cascade
openai-gpt-4.1-nano-llm
llama3.2-3b-llm
llama3.2-3b-cascade
llama3.2-1b-llm
glm-4.5-air-llm
gemma3-4b-llm
deepseek-flash-llm
deepseek-flash-cascade
0
2500
5000
7500
10000
12500
15000
17500
20000
warm latency, ms (bar = p95, tick = p50)