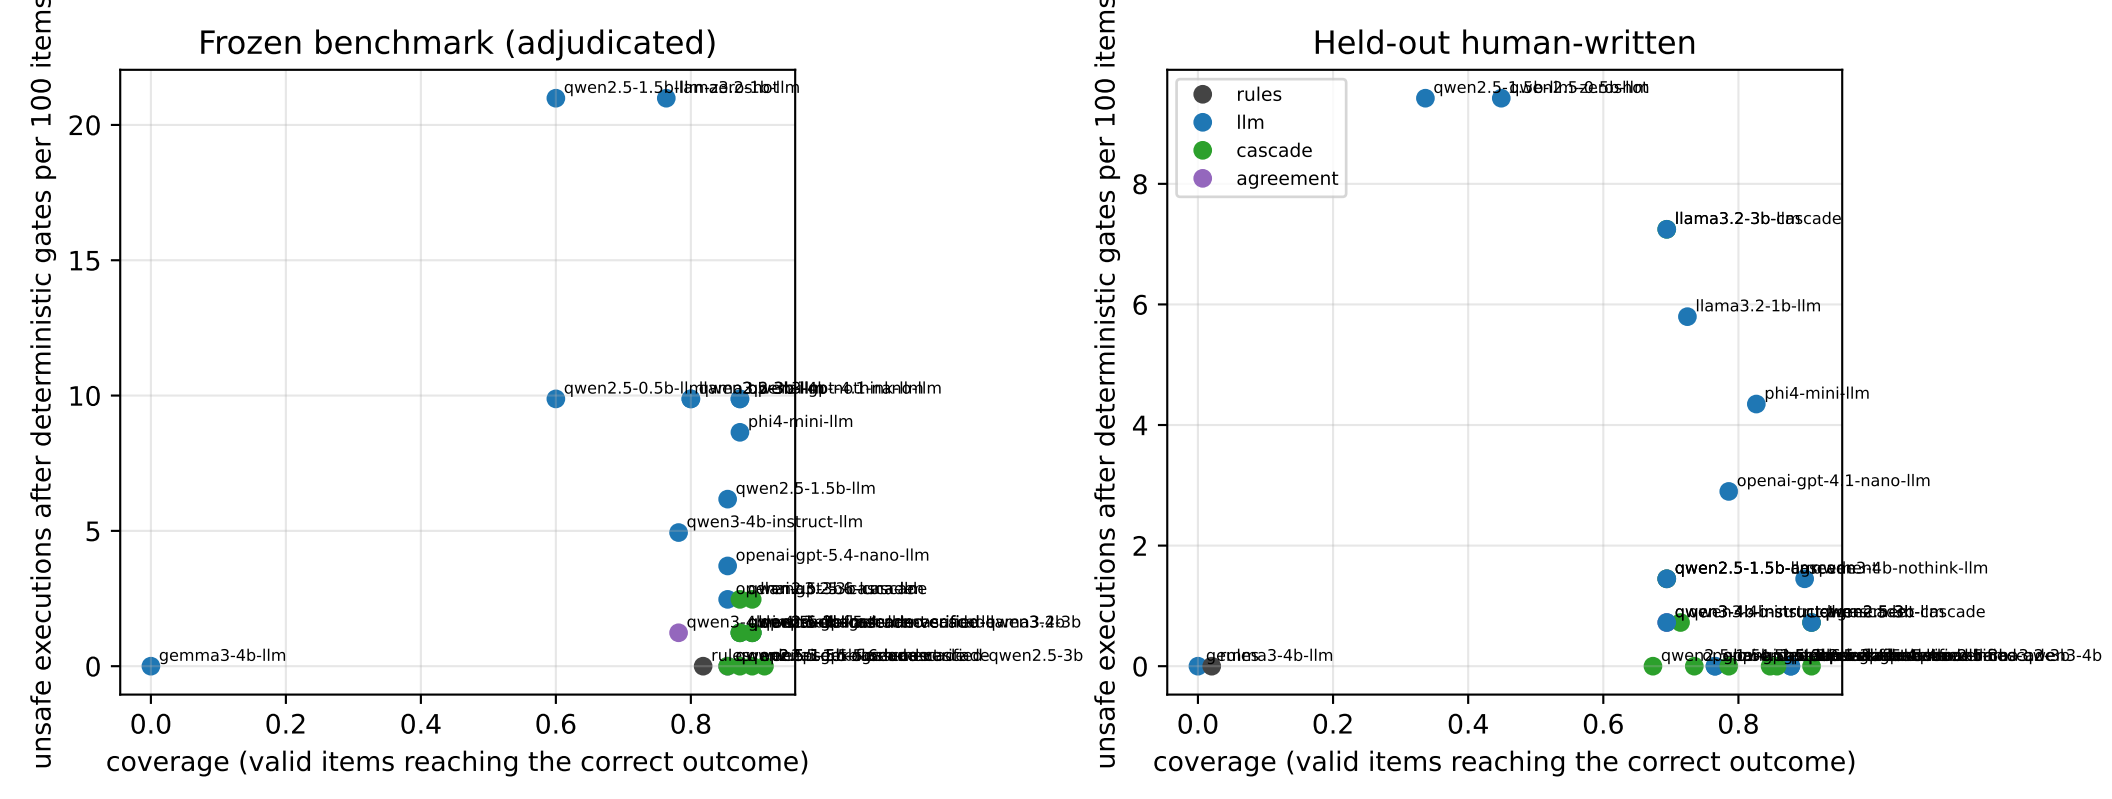
Frozen benchmark (adjudicated)
Held-out human-written
unsafe executions after deterministic gates per 100 items
coverage (valid items reaching the correct outcome)
rules
llm
cascade
agreement
qwen2.5-1.5b-llm
qwen2.5-0.5b-llm
phi4-mini-llm
qwen2.5-1.5b-llm
qwen3-4b-instruct-llm
openai-gpt-5.4-nano-llm
gemma3-4b-llm
llama3.2-3b-llm
llama3.2-1b-llm
phi4-mini-llm
openai-gpt-4.1-nano-llm
qwen2.5-1.5b-llm
gemma3-4b-llm
0.0
0.2
0.4
0.6
0.8
0
5
10
15
20
2
4
6
8